

## Formatting Submissions for a USENIX Conference: An (Incomplete) Example

Your N. Here  
*Your Institution*

Second Name  
*Second Institution*

## Abstract

Your abstract text goes here. Just a few facts. Whet our appetites. Not more than 200 words, if possible, and preferably closer to 150.

## 1 Introduction

[illegible]

## 2 Footnotes, Verbatim, and Citations

Footnotes should be places after punctuation characters, without any spaces between said characters and footnotes, like so.<sup>1</sup> And some embedded literal code may look as follows.

```
int main(int argc, char *argv[])
{
    return 0;
}
```

Now we're going to cite somebody. Watch for the cite tag. Here it comes. Arpachi-Dusseau and Arpachi-Dusseau co-authored an excellent OS Waldspurger got into the SIGOPS hall-of-fame due to his seminal paper

The tilde character (~) in the tex source means a non-breaking space. This way, your reference will always be attached to the word that preceded it, instead of going to the next line.

And the 'cite' package sorts your citations by their numerical order of the corresponding references at the end of the

---

<sup>1</sup>Remember that USENIX format stopped using endnotes and is now using regular footnotes.

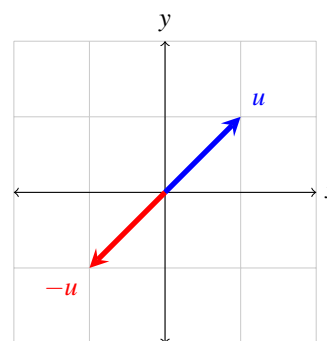


Figure 1: Text size inside figure should be as big as caption's text. Text size inside figure should be as big as caption's text. Text size inside figure should be as big as caption's text. Text size inside figure should be as big as caption's text. Text size inside figure should be as big as caption's text.

paper, ridding you from the need to notice that, e.g., “Waldspurger” appears after “Arpachi-Dusseau” when sorting references

It'd be nice and thoughtful of you to include a suitable link in each and every bibtex entry that you use in your submission, to allow reviewers (and other readers) to easily get to the cited work, as is done in all entries found in the References section of this document.

Now we're going to take a look at Section 3, but not before observing that refs to sections and citations and such are colored and clickable in the PDF because of the packages we've included.

### 3 Floating Figures and Lists

Here's a typical reference to a floating figure: that instructs you to make sure that the size of the text within the figures that you use is as big as (or bigger than) the size of the text in the caption of the figures. Please do. Really.

In our case, we've explicitly drawn the figure inlined in latex, to allow this tex file to cleanly compile. But usually, your figures will reside in some file.pdf, and you'd include them in your document with, say, `\includegraphics`.

Lists are sometimes quite handy. If you want to itemize things, feel free:

**fread** a function that reads from a `stream` into the array `ptr` at most `nobj` objects of size `size`, returning the number of objects read.

**Fred** a person's name, e.g., there once was a dude named Fred who separated `usenix.sty` from this file to allow for easy inclusion.

The noindent at the start of this paragraph in its tex version makes it clear that it's a continuation of the preceding paragraph, as opposed to a new paragraph in its own right.

### 3.1 LaTeX-ing Your TeX File

People often use `pdflatex` these days for creating pdf-s from tex files via the shell. And `bibtex`, of course. Works for us.

### Acknowledgments

The USENIX latex style is old and very tired, which is why there's no `acks` command for you to use when acknowledging. Sorry.

### Availability

### 3.2 Migration protocol

We observe the lifetime of a migratable data structure from the point it is created until after we finish migrating it to another node. Throughout the lifetime of the object and during the migration process, there are certain conditions that need to be satisfied by the application for the migration operation to finish correctly. As discussed later in this section, in many of these cases, forcing the application to satisfy low level conditions cannot be completely achieved by the library and the responsibility of conforming to these conditions will be left to the application.

of Slope at different stages during the above timeline.

#### 3.2.1 Object creation

During this phase the source machine calls into the Slope library multiple times, based on how many times memory allocation and deallocation is required.

**Source** creates a migratable object. Let us call this object the target object. Up until the point that the source initiates the migration, all of the memory allocations of the target object must happen through Slope's custom memory allocator through allocation contexts as discussed in The application must satisfy the memory allocation requirements discussed in

After the object is created the program will call the `run()` method of the object, which may result in further calls to the memory allocator as the object might need to allocate memory to process its incoming requests or carry out the calls to its member functions. These new allocations will be correctly picked up by Slope allocator given that allocation contexts are used know which memory pages belong to the target object.

#### 3.2.2 Migration initiation

The application logic decides that the target object must be migrated from the source machine to the destination. This might happen because the source machine is balancing out its load by offloading the object to the destination or because this particular object will benefit from running on the destination machine by being local to resources that are available there.

This is the first time that the destination machine will need to know about the properties of the target object. Had the source machine not initiated the migration, the destination machine would have stayed unaware of the existence of this object.

**Source** initiates the migration by calling into Slope. From this point on, the application instance on the source machine is neither allowed to cause any memory to be allocated to the target object, nor can it deallocate any memory that this object references. Each of these can result from calling member functions of the target object which either deallocate Slope memory previously held by the object, or create Slope memory allocation contexts from the object and use them to allocate Slope memory for the object.

This means the underlying Slope memory owned by this object is now fixed. The application can continue to read or write to the memory that the object already owns, until further on in the process when it receives a notification from the Slope library. This notification signals that source machine no longer owns the object, prohibiting this machine from any interaction with the target object, including reading from or writing to its memory. This will be discussed in detail in the following sections.

Slope posts a send request to the QP that is shared between the source and the destination, initiating the migration. With this request, the source will include  $n$  the total number of memory pages that the target object references. Note that Slope requires the memory pages owned by the target object to stay the same during the migration process after a call which initiates the migration, which means  $n$  will be constant throughout the process.

**Destination** receives a migration request along with  $n$ , the number of memory pages that the target object owns and therefore need to be transferred. The destination then populates another QP shared with the source with  $n$  receive requests, each corresponding to one of the virtual memory pages that underlie the target object. We switch to a different QP to allow concurrent migrations to happen between pairs of nodes. The destination then goes on to notify the source that it can receive the description of the  $n$  pages. The destination needs to know the page addresses to first create mappings for them and then pin them in physical memory so that their physical address will stay the same throughout the migration, while the network device writes to the pages by their physical address over DMA.

**Source** receives the clear to send from the destination and sends the starting address of each of the  $n$  pages to the destination.

**Destination** waits for  $n$  page addresses, and for each one of them, pins the corresponding page in physical memory, so that the addresses that underlie the target object can be used in RDMA read and write verbs. Notice how “pro-active” strategies which do not require knowledge of the page addresses will not work, as pinning the whole Slope memory in the physical memory will in the worst case, exhaust the available physical memory on some nodes, and in the best case, result in large amounts of wasted physical memory space. When the  $n$  pages are pinned and are ready to be the target of the RDMA verbs, the destination responds back, signaling that it has pinned all of the required pages.

### 3.2.3 Prefill

The goal of the prefill phase is to warm up the destination machine’s memory to minimize the window of time during which none of the two participating machines own the target object. No machine can read from or write to the object during that time. Therefore a long window without no owners means a long throughput and latency hiccup as the application does not make any progress.

During the prefill phase we copy the target object to the destination over RDMA. What makes it different from simply transferring the contents of memory is that during the prefill operation, the source is still allowed to write to any location in the target object memory, despite not having permission to change the memory layout of the object in any way by allocating or deallocating memory. We optimistically do the above transfer knowing that some of the pages might need to be retransferred as they are written to by the source machine after they are sent to the destination. We will refer to these pages as dirty pages. While pages can be dirtied, we go through the prefill phase, hoping that one of the machines (namely the source machine) will be able to partially function

during this phase, while dirtying a small percentage of the pages.

**Source** will need to loop through the memory pages of the target object to prepare them to be read safely by the network device. It will then need to use RDMA WRITE to send them to the destination, but before it is ready to go through these, it must put in place the means to detect the pages that are dirtied. A page will be dirtied when a write takes place in it, after it is sent to the destination.

**Dirty page detection:** This will be done using memory mapping protections and manually handling signals that are raised as a consequence of accesses to the protected memory pages.

First, the source overrides the default behavior of handling the SIGSEGV signal using a call to `sigaction`, to prevent the program from terminating on invalid memory references. It is important to keep in mind that SIGSEGV is delivered to the *same* thread whose current instruction reads or writes memory from a page that does not have the required permissions.

To synchronize the access to the pages of the target object, we set the protection flags of each page to `PROT_READ` before sending it over RDMA. This prevents any writes from happening while RDMA WRITE is in progress. These writes will result in a SIGSEGV signal being raised because of an invalid memory reference, and the thread will start executing our custom SIGSEGV handler.

Inside the signal handler, we make note of the address the access to which caused the signal to be raised. This can be found in the `si_addr` field from the `struct *siginfo_t` which is passed to the handler. Assuming this is an address in a page that the target object owns, we mark the page containing the address as dirty. In the simplest case, the RDMA WRITE corresponding to this page has previously finished. We update the protection flags of this page back to `PROT_READ | PROT_WRITE` to allow further writes to the addresses in this page to succeed.

Otherwise the RDMA WRITE is still in progress in which case we need to be more careful. In this case we need to wait for the RDMA WRITE to finish before adding the `PROT_WRITE` permission flag to the page. In our implementation we use a conditioned shared mutex where the thread which carries out the RDMA operations is the writer and the threads that may call into the signal handler are potential readers. As a result, each thread will block inside the signal handler until the address to which they are trying to write acquires the correct permission. As each of them they fall into one of the above cases while Slope prefills a page.

**Preparation:** At this point in the process, the source still has ownership over the target object memory. This means we need to coordinate the access to these pages to prevent

simultaneous reads and writes. To do this, the source uses `mprotect` to make the current page read-only. Semantically, the application is still allowed to write to the current page, except when the RDMA SEND is happening.

**Send:** Each page is posted to a data plane QP shared between the source and the sink, to be written to its corresponding virtual address in destination over RDMA. Notice that we do not put the `PROT_WRITE` permission of the current page back after the RDMA SEND finishes, as we need to detect any future writes which will dirty the page. However we do set a condition variable which a threads running the signal handler will check to know whether or not it is clear to put the `PROT_WRITE` flag back. At any point in time the clean pages will have their condition variable unset, which is how we distinguish between the dirty and clean pages.

One by one, the source prepares and writes each page that the target object owns, using the above procedures and keeps track of which of them is dirtied over time.

### 3.2.4 Transfer of ownership

Contrary to case of detecting the pages that become dirty, which can be done as described above, prevention of this phenomenon cannot be dealt with this way. For example we may try to set `PROT_NONE` permission on a page to prevent accesses to it after we have started to send it to the destination. The reason why this will not work will be discussed in the next section.

This cannot be correctly enforced by means such as setting the `PROT_NONE` on these pages. First, we might need multiple `mprotect` calls to accomplish this if the pages are not adjacent, meaning that we cannot assume the access to the object is ceased at a single point in time. Furthermore even if we are able to take away the access to all of the memory under the target object at once, if an application thread is currently in the process of writing to the above memory, there is no right action we can take from outside the application. For example this might have been one write operation in the middle of a series of writes that the application had been doing during the execution of a certain function. Therefore from the standpoint of the library, we do not have enough information to deal with cases similar to these, as aside from detecting the illegal writes, there is no right action to take to remedy them.

At this point we are ready to turn the ownership over to the destination. We also have to notify the application at the source to stop referencing the target object for both reads and writes before we can proceed.

**Destination** receives the completions for the prefill RDMA WRITES. That is because those WRITES are done with immediate values which will results in completion entries to be created for them. After the destination receives the completion of the last SEND, it will send a request back to the

source, asking the source to finally give up the ownership of the object.

**Source** calls back to the application upon receiving the above message to make sure the application will not reference the target object for either of reads or writes after the callback returns. The application is therefore responsible to block the callback and notifies all of its threads to finish their ongoing operations on the target object. The application will ideally transition to a “pre-transfer” mode after the initial migration initiation. In this phase the application will try to do mostly read-only accesses to avoid dirtying too many pages as this will cause the migration to be done less efficiently with more pages that need to be transferred twice. discuss how this model fits naturally in C++ programs.

After the application clears the migration, effectively promising to never reference the target object again, Slope will use the dirty page tracking data to identify the dirty pages and send their addresses to the destination. This message will also hand over the ownership of the target object to the destination.

**Destination** is notified that it has been given the ownership of the object and is clear to start the `run()` method of the object, but also needs to re-pull the dirty pages whose address is attached to the request since the contents of these pages on the destination machine are outdated. Slope starts the above two processes concurrently. The destination machine first sets `PROT_READ | PROT_WRITE` access for the clean pages and `PROT_NONE` for the dirty pages. We put the dirty page addresses into a priority queue and fetch them from the source machine in order, giving precedence to the pages which result in a `SIGSEGV` being raised at the destination from an access by the application. This means that the page has `PROT_NONE` permission set and is immediately required by an application thread. When the page is pulled using RDMA READ, the page mapping flags will be updated to include read and write permissions and any thread waiting for those pages can proceed. This is done similar to how we do the dirty page detection on the source machine which was discussed earlier.

### 3.2.5 Post-transfer

After the transfer is done, the state of the application will be indistinguishable from a situation where we allocated the target object on the destination machine from the start, if we deal with one final housekeeping task.

**Destination** will notify the source that it has READ all of the dirty pages and that it no longer needs the destination to keep the pre-transfer content of the object memory.

**Source** will wait to receive the above message at which point it safely `munmaps` the memory underlying the “old” target object. At this point the state on the source looks as if the said object had never existed there.

#### **4 Conformance to C++ language terminologies**

e.g. `const`, `mutable`, `const cast`

#### **5 Optimizations and Corner cases**

e.g. have to send the location of the object itself also: from which address to create the migration pointer.

USENIX program committees give extra points to submissions that are backed by artifacts that are publicly available. If you made your code or data available, it’s worth mentioning this fact in a dedicated section.